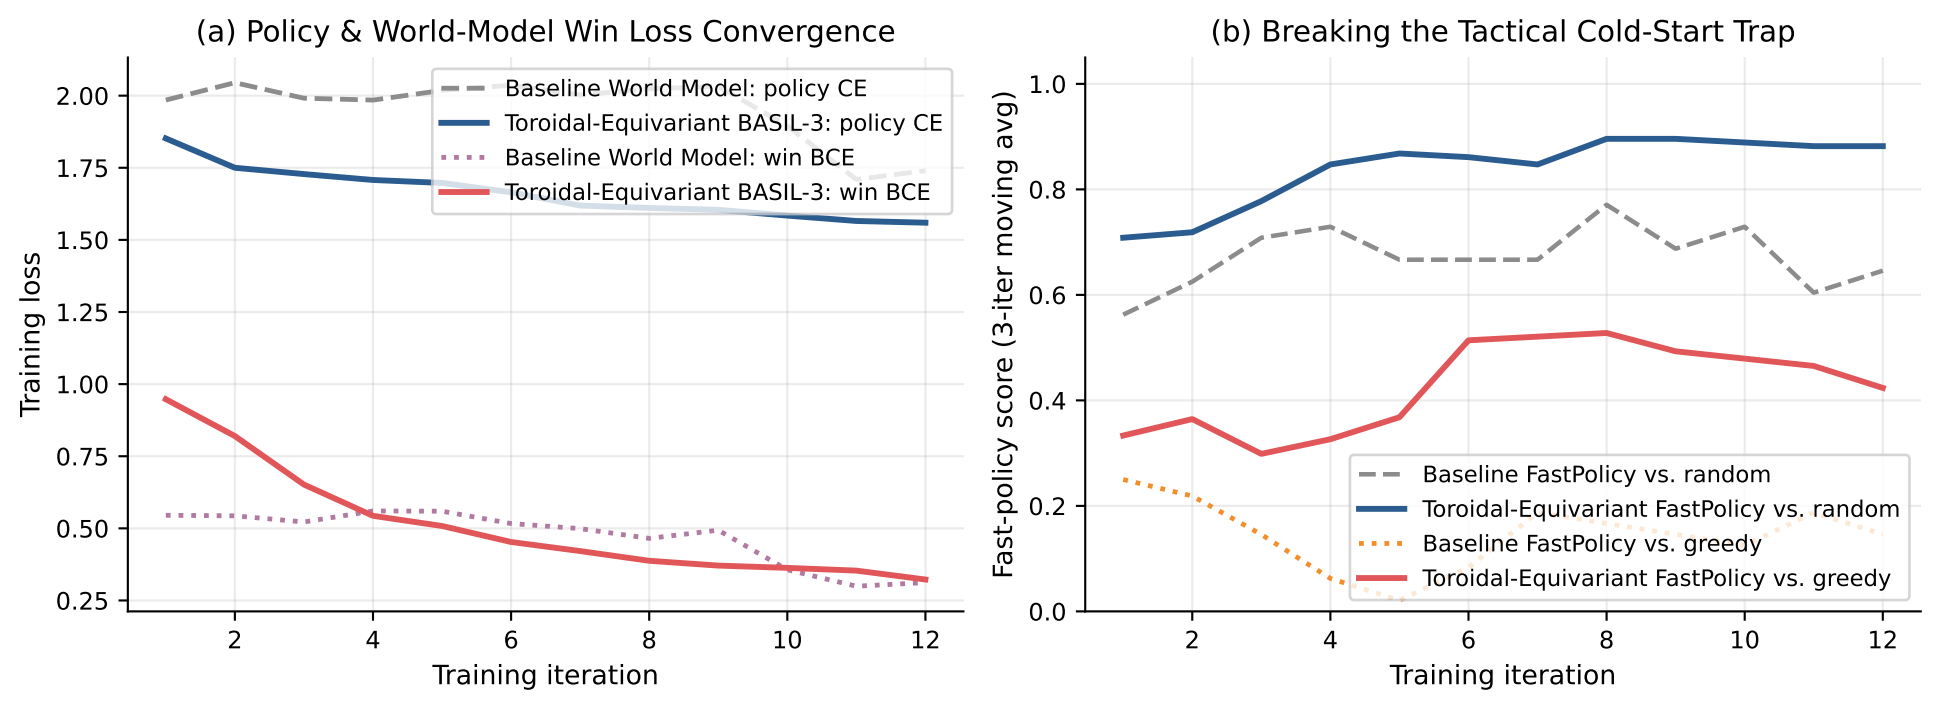
(a) Policy & World-Model Win Loss Convergence
Baseline World Model: policy CE
Toroidal-Equivariant BASIL-3: policy CE
Baseline World Model: win BCE
Toroidal-Equivariant BASIL-3: win BCE
Training loss
Training iteration
(b) Breaking the Tactical Cold-Start Trap
Baseline FastPolicy vs. random
Toroidal-Equivariant FastPolicy vs. random
Baseline FastPolicy vs. greedy
Toroidal-Equivariant FastPolicy vs. greedy
Fast-policy score (3-iter moving avg)
Training iteration