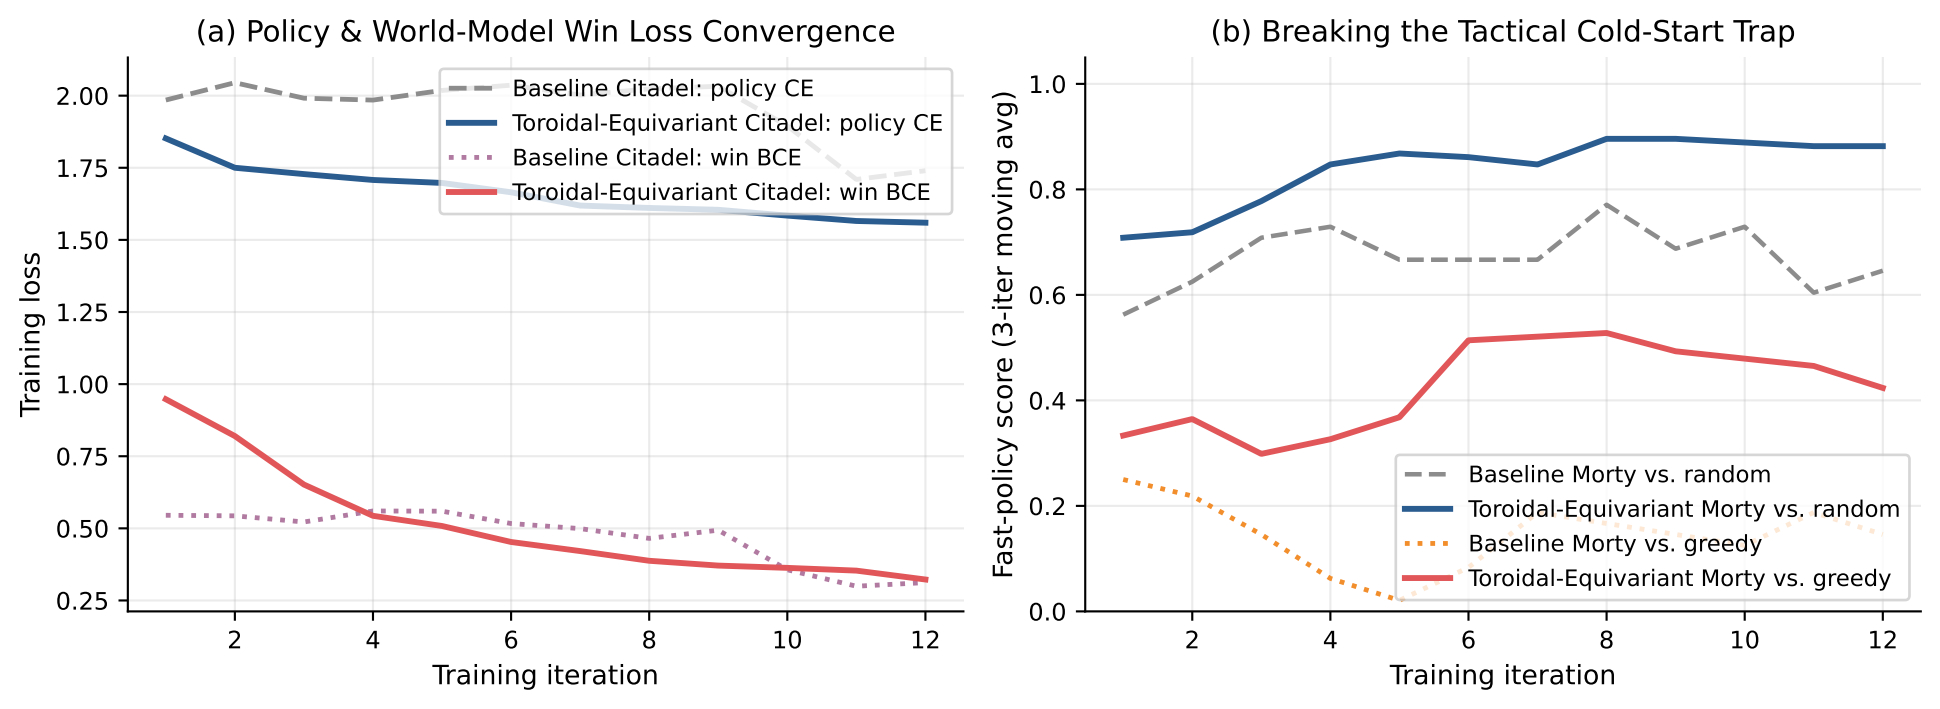
(a) Policy & World-Model Win Loss Convergence
Training loss
Training iteration
Baseline Citadel: policy CE
Toroidal-Equivariant Citadel: policy CE
Baseline Citadel: win BCE
Toroidal-Equivariant Citadel: win BCE
(b) Breaking the Tactical Cold-Start Trap
Fast-policy score (3-iter moving avg)
Training iteration
Baseline Morty vs. random
Toroidal-Equivariant Morty vs. random
Baseline Morty vs. greedy
Toroidal-Equivariant Morty vs. greedy